Data basics

ME 447/547 Visualizing Data

Richard Layton

March 2019

Rose-Hulman Institute of Technology

Preparing data for graphs starts with four basic skills

Obtain the raw data

Read raw data into R and examine it

Identify the structure of your data

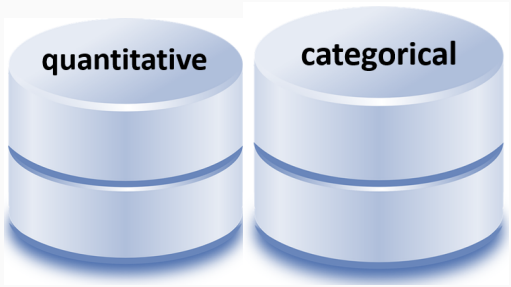

Tidy the data and write to file

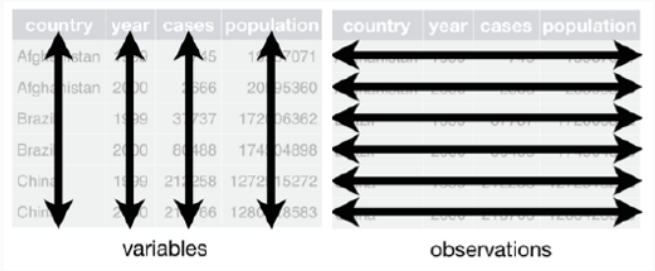

Data are everywhere

Data are provided in base R

Data are provided in R packages

Online sources are ubiquitous

- FiveThirtyEight `https://data.fivethirtyeight.com/`
- US government `https://www.data.gov/`
- NOAA climate data `https://www.ncdc.noaa.gov/cdo-web/`
- Publications for which code and/or data are available

  `https://reproducibleresearch.net/reproducible-material/`

You may even have data of your own from prior courses or research

For practice, use data in base R

`data()` to list data sets in base R

```
#> AirPassengers  Monthly Airline Passenger Numbers
#> BJsales        Sales Data with Leading Indicator
#> BOD            Biochemical Oxygen Demand
#> CO2            Carbon Dioxide Uptake in Grass Plants
#> Formaldehyde   Determination of Formaldehyde
etc.
```

For practice, use data in R packages

`data(package = "dplyr")` to list data sets in package dplyr

```
#> band_instruments   Band membership
#> band_instruments2  Band membership
#> band_members       Band membership
#> nasa               NASA spatio-temporal data
#> starwars           Starwars characters
#> storms             Storm tracks data
```

? data_name shows the data set help page

```
library("graphclassmate")
data(package = "graphclassmate")
? metro_pop
```

metro_pop {graphclassmate} R Documentation

Population in the NY metro area

Description

A data set of population in the New York metropolitan area by county and race/ethinicty from the 2000 census.

Usage

```
metro_pop
```

Format

A tidy data frame (tibble) with 60 observations and 3 variables. An observation is the population in a county by race/ethnicity.

race

Race or ethnicity

county

Name of county

population

Number of residents from the 2000 US census

Data in base R and in R packages are automatically loaded

Launching R loads all data sets in base-R

Loading a package with `library()` loads all the data sets in the package

All other data files have to be read or web-scraped

Save original data files in the data-raw directory

Data in their original form are never edited manually

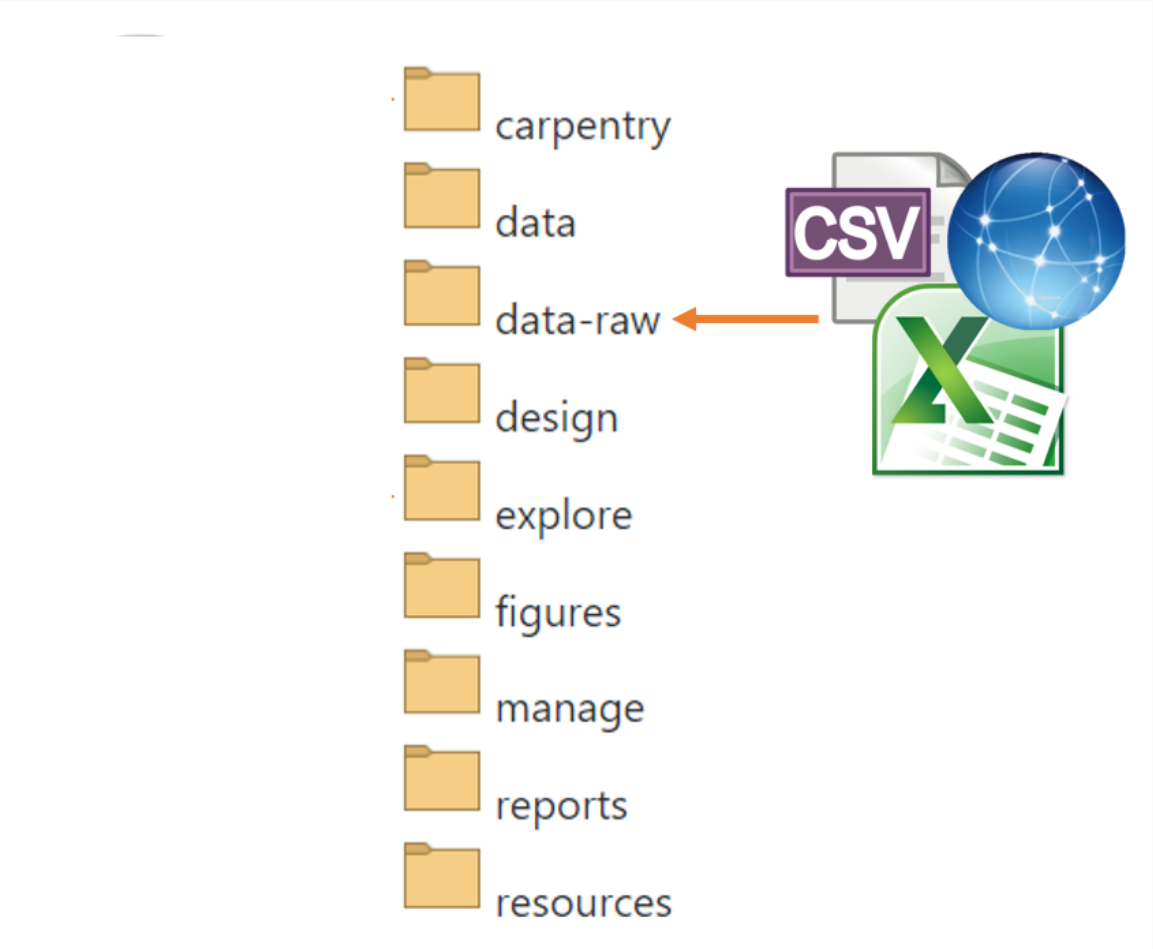

We work with file manangement in detail during the data studio.

Read the raw data with R scripts

R scripts are saved in the carpentry or explore directories

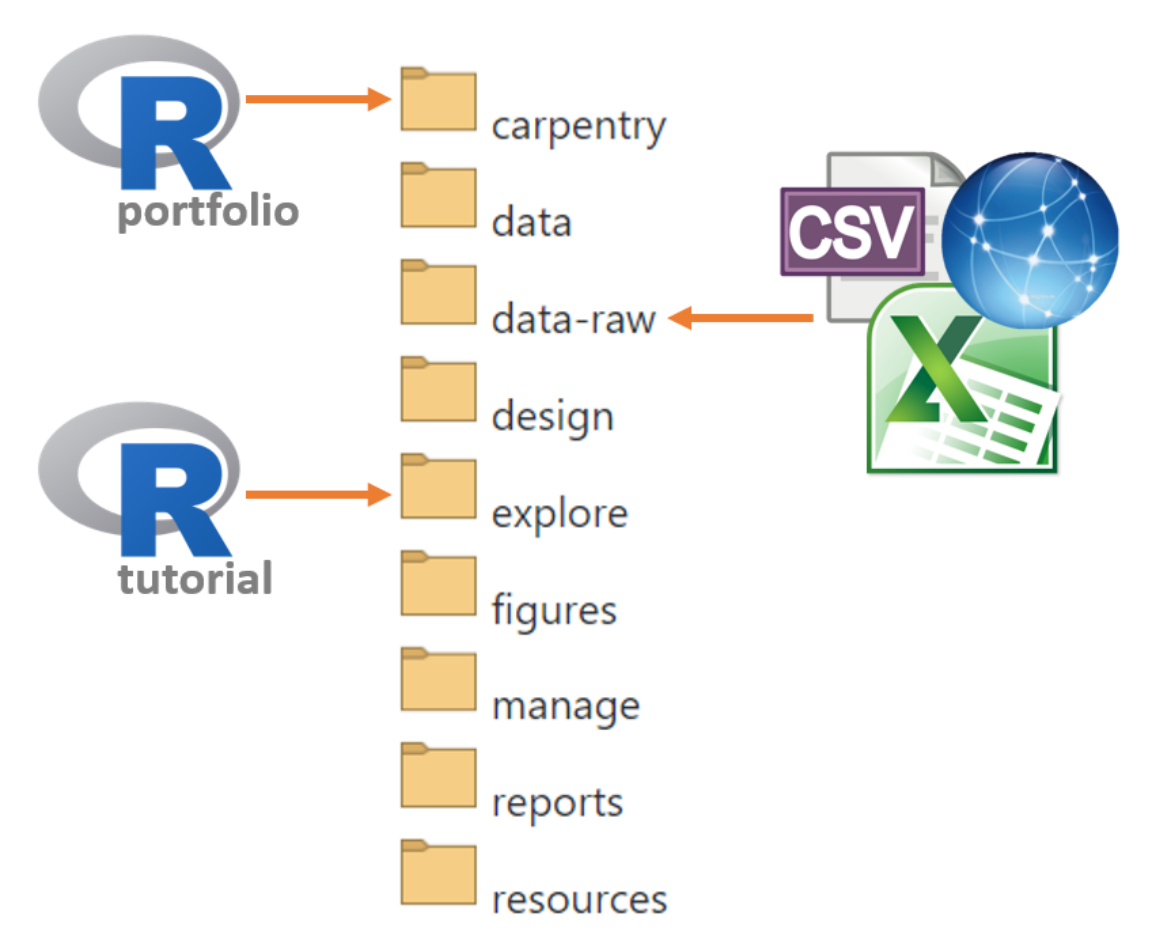

We work with file manangement in detail during the data studio.

Suppose `data-raw/` contains data in an Excel file

readxl is the package (you will have to install the package)

`read_excel()` is the function

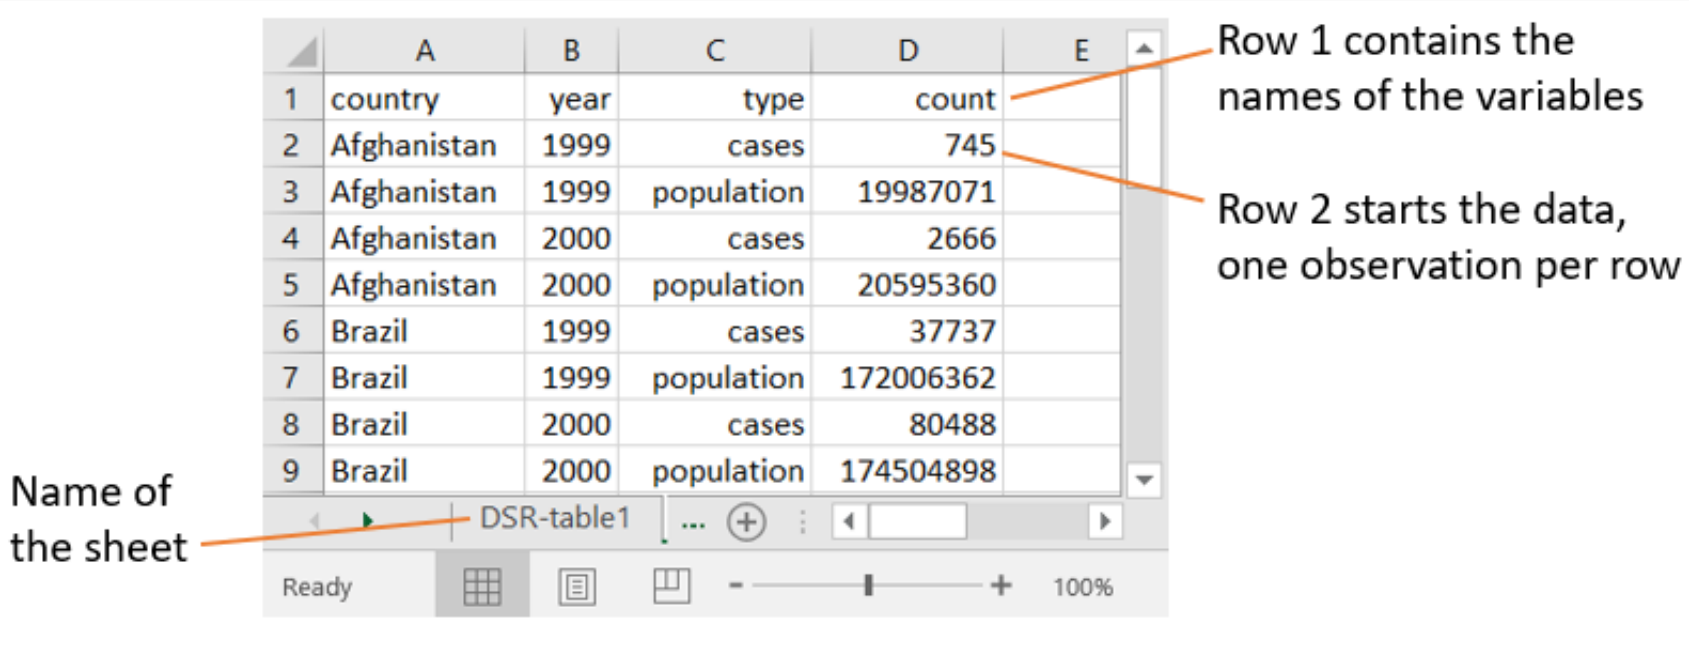

	A	B	C	D	E
1	country	year	type	count	
2	Afghanistan	1999	cases	745	
3	Afghanistan	1999	population	19987071	
4	Afghanistan	2000	cases	2666	
5	Afghanistan	2000	population	20595360	
6	Brazil	1999	cases	37737	
7	Brazil	1999	population	172006362	
8	Brazil	2000	cases	80488	
9	Brazil	2000	population	174504898	

We work with data tidying in detail during the data studio.

`read_excel()` to read an Excel file

```
library("readxl")
tidy_data <- read_excel(path  = "data-raw/DSR-table1.xlsx",
                        sheet = "DSR-table1")
```

We can pretty-print the data using `knitr::kable()`

```
library("knitr")
kable(tidy_data)
```

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583

Suppose `data-raw/` contains data in a CSV file

readr is the package (part of the tidyverse)

`read_csv()` is the function

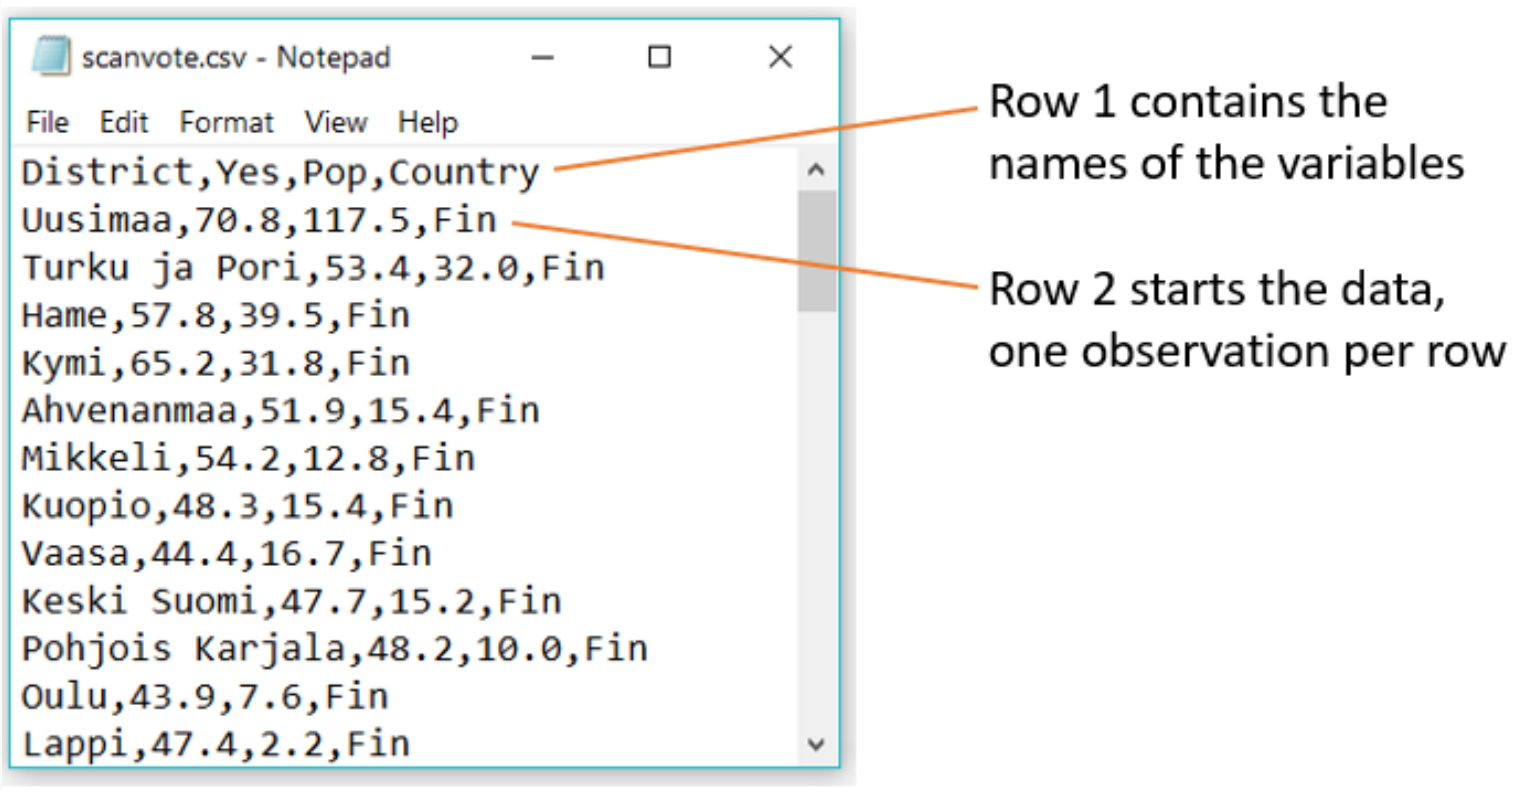

We work with data tidying in detail during the data studio.

`read_csv()` to read a CSV file

```
library("tidyverse") # loads the readr package
tidy_data_2 <- read_csv(file = "data-raw/scanvote.csv")
```

We can pretty-print the top n rows with `head()`

```
tidy_data_2 %>%
    head(., n = 5L) %>%
    kable()
```

District	Yes	Pop	Country
Uusimaa	70.8	117.5	Fin
Turku ja Pori	53.4	32.0	Fin
Hame	57.8	39.5	Fin
Kymi	65.2	31.8	Fin
Ahvenanmaa	51.9	15.4	Fin

`read_excel()` and `read_csv()` produce tibbles

```
class(tidy_data)
```

```
#> [1] "tbl_df"     "tbl"        "data.frame"
```

```
class(tidy_data_2)
```

```
#> [1] "spec_tbl_df" "tbl_df"      "tbl"         "data.frame"
```

Confine your webscraping (for now) to data in ASCII format

Introduction
Methodology
Data Explanation

CANADA

- Newfoundland
- Prince Edward Island
- Nova Scotia
- New Brunswick
- Quebec
- Ontario
- Manitoba
- Saskatchewan
- Alberta
- British Columbia
- Northwest Territories
- Yukon

Archives (Life Tables) Links

Canadian Human Mortality Database

Alberta

CHMD Data 1921-2011	Age Interval x Year Interval					
	1 x 1	1 x 5	1 x 10	5 x 1	5 x 5	5 x 10
Births	View					
Deaths (Lexis triangle)	View			View		
Population size	View			View		
Exposure-to-risk	View	View	View	View	View	View
Death rates	View	View	View	View	View	View
Life tables - Male	View	View	View	View	View	View
Life tables - Female	View	View	View	View	View	View
Life tables - Total	View	View	View	View	View	View
Life expectancy at birth	View					

Source: Canadian Human Mortality Database

On any online data page, right-click > View page source

Canada-Alberta, Population size (1-year) Last modified: 31-Jul-2014, MPv5 (May07)

Year	Age	Female	Male	Total
1921	0	7864.85	8133.86	15998.71
1921	1	7936.45	8142.91	16079.36
1921	2	8024.81	8240.34	16265.15
1921	3	8017.72	8244.01	16261.73
1921	4	7925.59	8154.78	16080.37
1921	5	7760.15	7991.74	15751.89
1921	6	7530.77	7768.78	15299.55
1921	7	7250.42	7495.09	14745.51
1921	8	6926.93	7181.14	14108.07
1921	9	6568.14	6829.25	13397.39
1921	10	6177.87	6440.22	12618.09
1921	11	5832.52	6104.80	11937.32
1921	12	5561.03	5857.21	11418.24
1921	13	5341.23	5669.97	11011.20
1921	14	5112.95	5478.27	10591.22
1921	15	4892.53	5302.30	10194.83
1921	16	4701.44	5147.45	9848.89
1921	17	4557.01	5008.00	9565.01
1921	18	4443.74	4893.96	9337.70
1921	19	4335.37	4802.62	9137.99
1921	20	4242.44	4727.01	8969.45

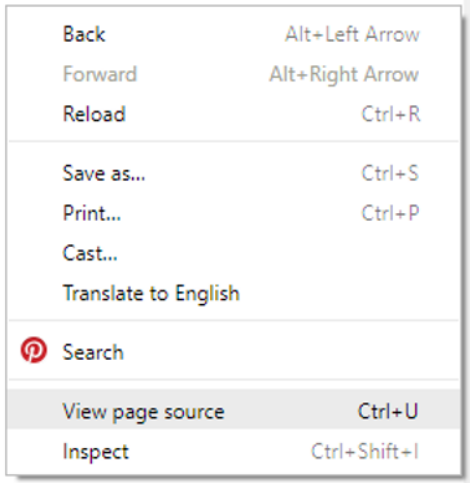

Data formatted in ASCII (text) is easily recognized

Canada-Alberta, Population size (1-year) Last modified: 31-Jul-2014, MPv5 (May07)

Year	Age	Female	Male	Total
1921	0	7864.85	8133.86	15998.71
1921	1	7936.45	8142.91	16079.36
1921	2	8024.81	8240.34	16265.15
1921	3	8017.72	8244.01	16261.73
1921	4	7925.59	8154.78	16080.37
1921	5	7760.15	7991.74	15751.89
1921	6	7530.77	7768.78	15299.55
1921	7	7250.42	7495.09	14745.51
1921	8	6926.93	7181.14	14108.07
1921	9	6568.14	6829.25	13397.39
1921	10	6177.87	6440.22	12618.09
1921	11	5832.52	6104.80	11937.32
1921	12	5561.03	5857.21	11418.24
1921	13	5341.23	5669.97	11011.20
1921	14	5112.95	5478.27	10591.22
1921	15	4892.53	5302.30	10194.83
1921	16	4701.44	5147.45	9848.89
1921	17	4557.01	5008.00	9565.01

Data formatted in HTML is also easily recognized

```
<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 4.0 TRANSITIONAL//EN"
 "http://www.w3.org/TR/REC-html40/loose.dtd">
<HTML LANG="en">
<head>
<script>
if (document.layers)
  WM_scaleFont(initialFontSize, fontUnits);
</script>
<title>Historical Census of Housing Tables Home Values - Housing Topics - U.S. Census Bureau</TITLE>

<meta http-equiv="Content-Type" content="text/html; charset=iso-8859-1" />

<meta name="DC.title" content="US Census Bureau Historical Census of Housing Tables Home Values" />

<meta name="DC.description" content="Selected housing characteristics data from decennial census housing files are presented here for the United States and for each state. Trend analyses are discussed, with graphic illustration at the national level." />

<meta name="DC.creator" content="SEHSD Division" />

<meta name="DC.date.created" scheme="ISO8601" content="2000-06-01" />

<meta name="DC.date.reviewed" scheme="ISO8601" content="2000-06-01" />

<meta name="DC.language" scheme="DCTERMS.RFC1766" content="EN-US" />

<meta name="author" content="US Census Bureau Historical Census of Housing Tables Home
```

Source: US Census of Housing

With online data in ASCII format, webscraping is easy

`utils` is the package

`read./table()` is the function

```
library("utils")
url <-
  "http://www.prdh.umontreal.ca/BDLC/data/alb/Population.txt"

df  <- read.table(url,
                  skip = 2,
                  header = TRUE,
                  stringsAsFactors = FALSE)

df <- as_tibble(df)
```

Examine it and write it to the `data-raw` directory

```
glimpse(df)
#> Observations: 10,212
#> Variables: 5
#> $ Year   <int> 1921, 1921, 1921, 1921, 1921, 1921, 192
#> $ Age    <chr> "0", "1", "2", "3", "4", "5", "6", "7",
#> $ Female <dbl> 7864.85, 7936.45, 8024.81, 8017.72, 792
#> $ Male   <dbl> 8133.86, 8142.91, 8240.34, 8244.01, 815
#> $ Total  <dbl> 15998.71, 16079.36, 16265.15, 16261.73,

write_csv(df, "data-raw/alberta_mortality.csv")
```

When the data are not tidy, ...

VADeaths.xlsx

	A	B	C	D	E
1		Rural		Urban	
2	Group	Men	Women	Men	Women
3	50-54	11.7	8.7	15.4	8.4
4	55-59	18.1	11.7	24.3	13.6
5	60-64	26.9	20.3	37	19.3
6	65-69	41	30.9	54.6	35.1
7	70-74	66	54.3	71.1	50

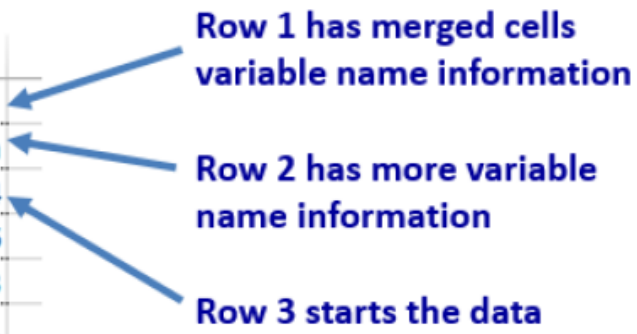

… the read results can be weird.

```
untidy_data <- read_excel(path = "data-raw/VADeaths.xlsx",
                          sheet = "VADeaths") %>%
               glimpse()
```

```
#> Observations: 6
#> Variables: 5
#> $ `..1` <chr> "Group", "50-54", "55-59", "60-64", "65-69", "
#> $ Rural <chr> "Men", "11.7", "18.100000000000001", "26.9", "
#> $ `..3` <chr> "Women", "8.6999999999999993", "11.7", "20.3",
#> $ Urban <chr> "Men", "15.4", "24.3", "37", "54.6", "71.09999
#> $ `..5` <chr> "Women", "8.4", "13.6", "19.3", "35.1", "50"
```

All the cells have been converted to character data

The result is more easily seen using `knitr::kable()`

```
kable(untidy_data)
```

..1	Rural	..3	Urban	..5
Group	Men	Women	Men	Wom
50-54	11.7	8.699999999999993	15.4	8.4
55-59	18.100000000000001	11.7	24.3	13.6
60-64	26.9	20.3	37	19.3
65-69	41	30.9	54.6	35.1
70-74	66	54.3	71.099999999999994	50

The first row is not an observation—that's the problem

When reading the file, we need to skip the first row

The data are at least readable but we have lost information

```
untidy_data <- read_excel(path = "data-raw/VADeaths.xlsx",
                          sheet = "VADeaths",
                          skip = 1)
```

Group	Men..2	Women..3	Men..4	Women..5
50-54	11.7	8.7	15.4	8.4
55-59	18.1	11.7	24.3	13.6
60-64	26.9	20.3	37.0	19.3
65-69	41.0	30.9	54.6	35.1
70-74	66.0	54.3	71.1	50.0

Carpentry on untidy data is a large fraction of your effort

It is often said that 80% of data analysis is spent on the process of cleaning and preparing the data.

Data preparation is not just a first step, but must be repeated many times over the course of analysis as new problems come to light or new data is collected.

—Hadley Wickham, Tidy Data

Understanding data structure is necessary for tidying

Number of variables?

Continuous or discrete?

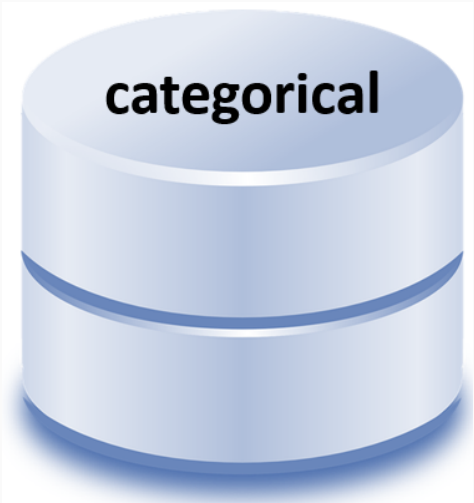

Number of variables?

Nominal or ordinal?

Number of levels each?

Transforming data to tidy form is necesaary for productivity

- how you tidy the data set before graphing
- the graph types that are suitable
- how easy it is to get ggplot to do your bidding
- how productively you spend your time

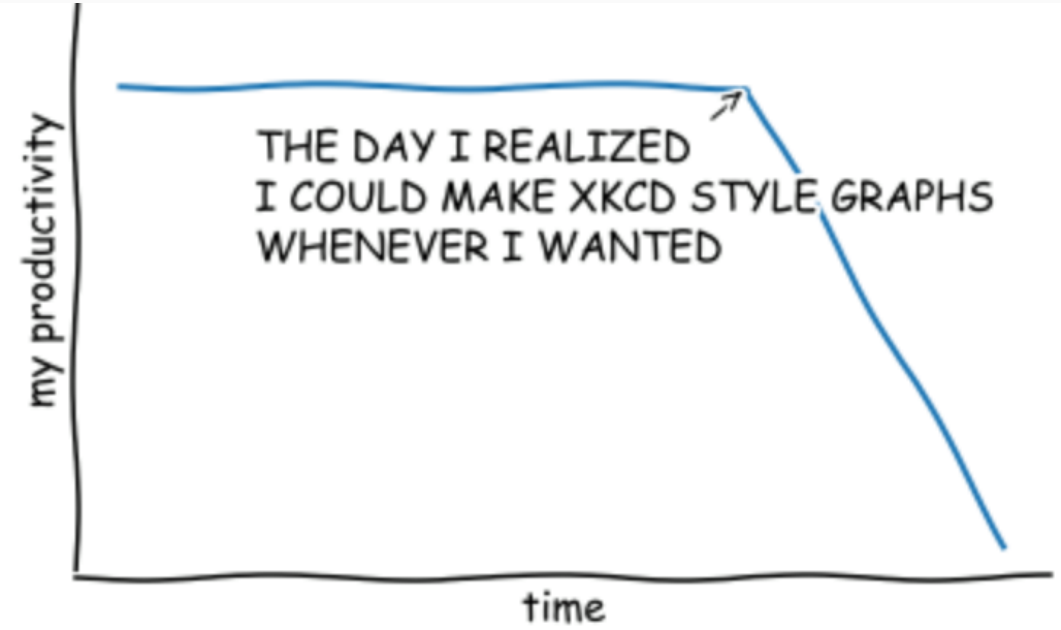

For graphical productivity, data has to be tidy

In a tidy data set:

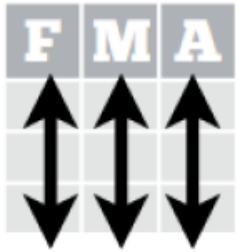

&

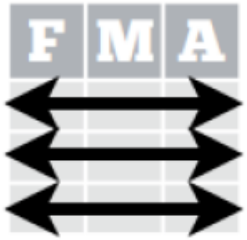

Each **variable** is saved in its own **column**

Each **observation** is saved in its own **row**

Source: data-wrangling-cheatsheet, https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf

Every column is a variable, every row is an observation

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583

table1

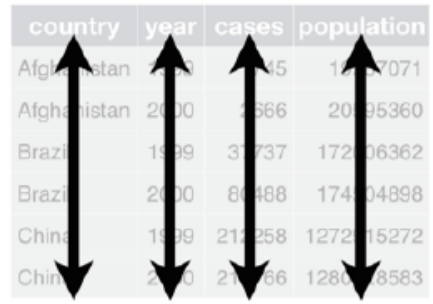

Source: Data Science with R by Garrett Grolemund, http://garrettgman.github.io/tidying/

Sadly, untidy data is common

country	year	key	value
Afghanistan	1999	cases	745
Afghanistan	1999	population	19987071
Afghanistan	2000	cases	2666
Afghanistan	2000	population	20595360
Brazil	1999	cases	37737
Brazil	1999	population	172006362
Brazil	2000	cases	80488
Brazil	2000	population	174504898
China	1999	cases	212258
China	1999	population	1272915272
China	2000	cases	213766
China	2000	population	1280428583

table2

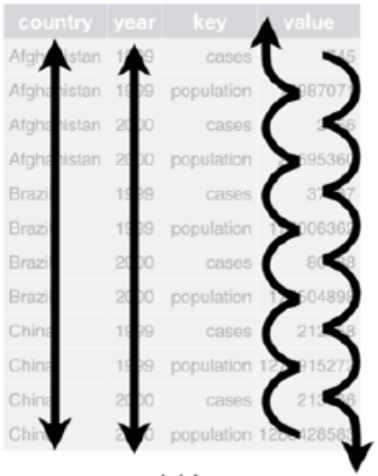

variables

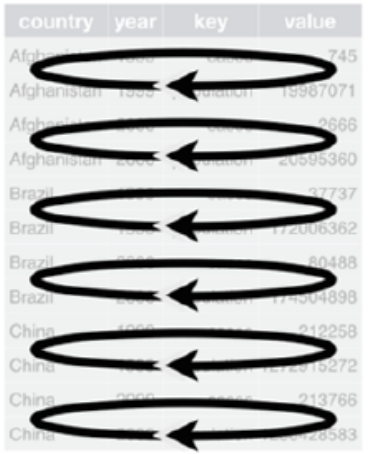

observations

Sadly, untidy data is common

country	1999	2000
Afghanistan	745	2666
Brazil	37737	80488
China	212258	213766

table4

country	1999	2000
Afghanistan	19987071	20595360
Brazil	172006362	174504898
China	1272915272	1280428583

table5

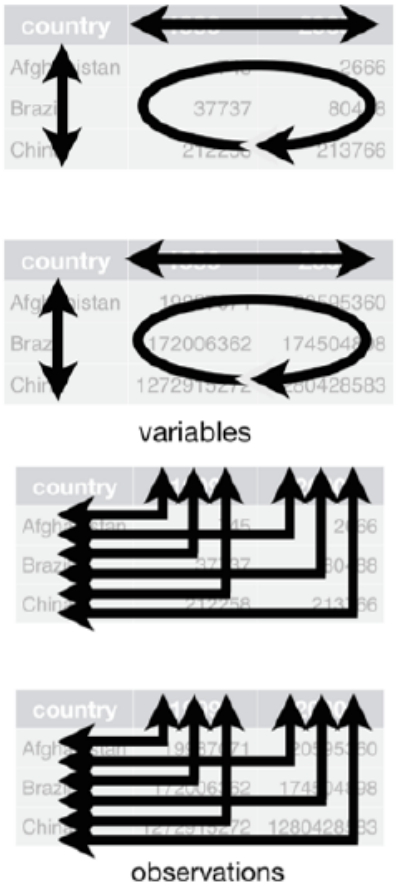

Source: Data Science with R by Garrett Grolemund, http://garrettgman.github.io/tidying/

Some industry or government spreadsheets are horribly untidy

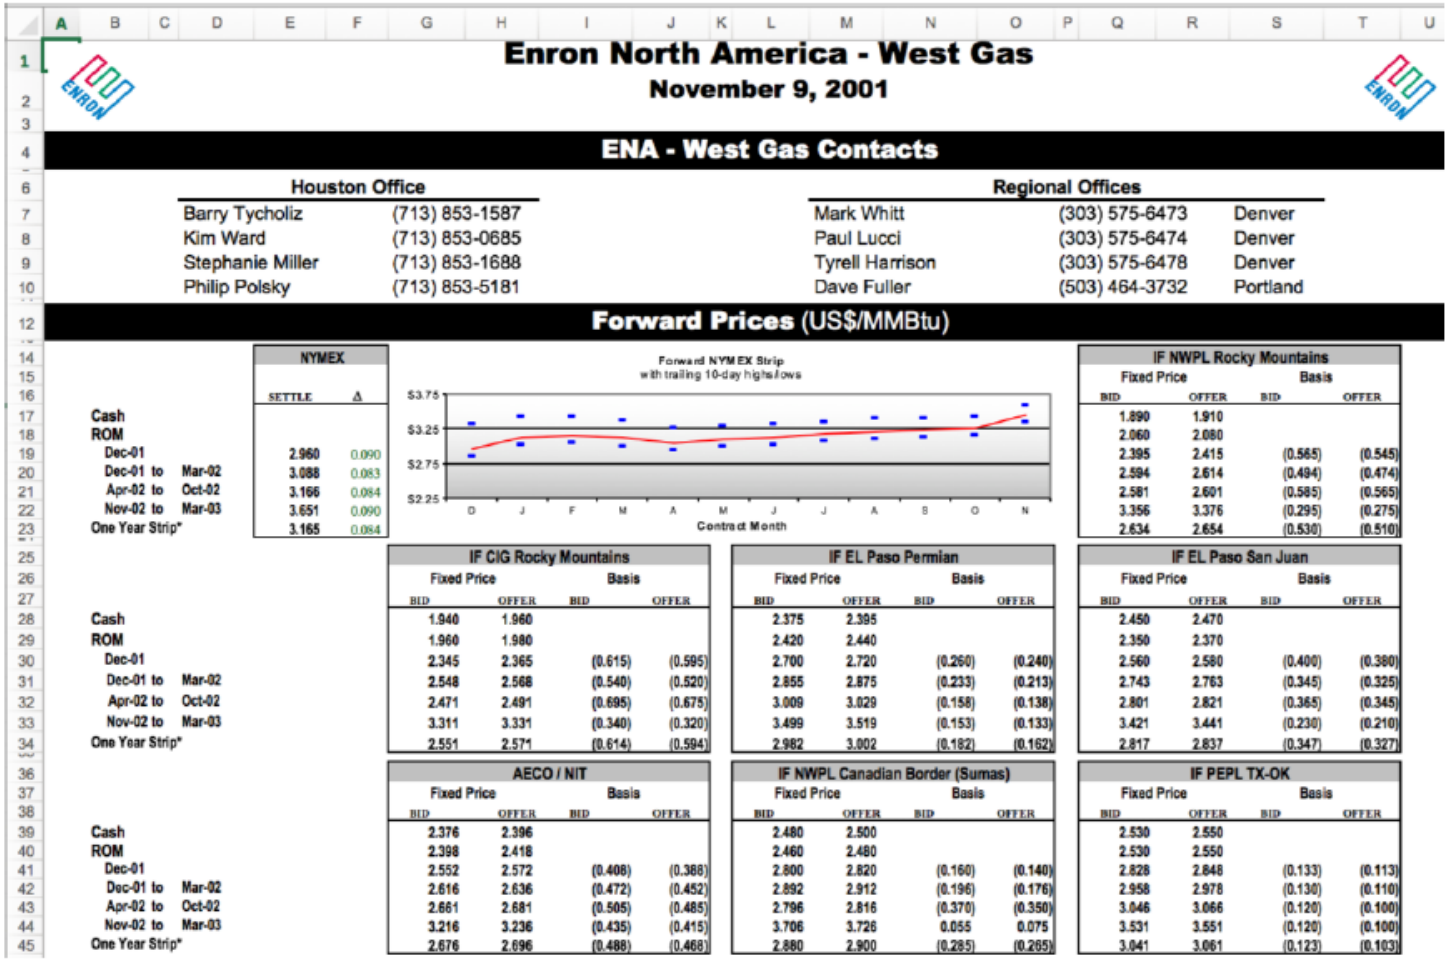

Enron North America - West Gas

November 9, 2001

ENA - West Gas Contacts

Houston Office		Regional Offices		
Barry Tycholiz	(713) 853-1587	Mark Whitt	(303) 575-6473	Denver
Kim Ward	(713) 853-0685	Paul Lucci	(303) 575-6474	Denver
Stephanie Miller	(713) 853-1688	Tyrell Harrison	(303) 575-6478	Denver
Philip Polsky	(713) 853-5181	Dave Fuller	(503) 464-3732	Portland

Forward Prices (US$/MMBtu)

	NYMEX SETTLE	NYMEX Δ
Cash		
ROM		
Dec-01	2.960	0.090
Dec-01 to Mar-02	3.088	0.083
Apr-02 to Oct-02	3.166	0.084
Nov-02 to Mar-03	3.651	0.090
One Year Strip*	3.165	0.084

Forward NYMEX Strip with trailing 10-day highs/lows

	IF NWPL Rocky Mountains Fixed Price BID	Fixed Price OFFER	Basis BID	Basis OFFER
Cash	1.890	1.910		
ROM	2.060	2.080		
Dec-01	2.395	2.415	(0.565)	(0.545)
Dec-01 to Mar-02	2.594	2.614	(0.494)	(0.474)
Apr-02 to Oct-02	2.581	2.601	(0.585)	(0.565)
Nov-02 to Mar-03	3.356	3.376	(0.295)	(0.275)
One Year Strip*	2.634	2.654	(0.530)	(0.510)

	IF CIG Rocky Mountains Fixed Price BID	Fixed Price OFFER	Basis BID	Basis OFFER
Cash	1.940	1.960		
ROM	1.960	1.980		
Dec-01	2.345	2.365	(0.615)	(0.595)
Dec-01 to Mar-02	2.548	2.568	(0.540)	(0.520)
Apr-02 to Oct-02	2.471	2.491	(0.695)	(0.675)
Nov-02 to Mar-03	3.311	3.331	(0.340)	(0.320)
One Year Strip*	2.551	2.571	(0.614)	(0.594)

	IF EL Paso Permian Fixed Price BID	Fixed Price OFFER	Basis BID	Basis OFFER
Cash	2.375	2.395		
ROM	2.420	2.440		
Dec-01	2.700	2.720	(0.260)	(0.240)
Dec-01 to Mar-02	2.855	2.875	(0.233)	(0.213)
Apr-02 to Oct-02	3.009	3.029	(0.158)	(0.138)
Nov-02 to Mar-03	3.499	3.519	(0.153)	(0.133)
One Year Strip*	2.982	3.002	(0.182)	(0.162)

	IF EL Paso San Juan Fixed Price BID	Fixed Price OFFER	Basis BID	Basis OFFER
Cash	2.450	2.470		
ROM	2.350	2.370		
Dec-01	2.560	2.580	(0.400)	(0.380)
Dec-01 to Mar-02	2.743	2.763	(0.345)	(0.325)
Apr-02 to Oct-02	2.801	2.821	(0.365)	(0.345)
Nov-02 to Mar-03	3.421	3.441	(0.230)	(0.210)
One Year Strip*	2.817	2.837	(0.347)	(0.327)

	AECO / NIT Fixed Price BID	Fixed Price OFFER	Basis BID	Basis OFFER
Cash	2.376	2.396		
ROM	2.398	2.418		
Dec-01	2.552	2.572	(0.408)	(0.388)
Dec-01 to Mar-02	2.616	2.636	(0.472)	(0.452)
Apr-02 to Oct-02	2.661	2.681	(0.505)	(0.485)
Nov-02 to Mar-03	3.216	3.236	(0.435)	(0.415)
One Year Strip*	2.676	2.696	(0.488)	(0.468)

	IF NWPL Canadian Border (Sumas) Fixed Price BID	Fixed Price OFFER	Basis BID	Basis OFFER
Cash	2.480	2.500		
ROM	2.460	2.480		
Dec-01	2.800	2.820	(0.160)	(0.140)
Dec-01 to Mar-02	2.892	2.912	(0.196)	(0.176)
Apr-02 to Oct-02	2.796	2.816	(0.370)	(0.350)
Nov-02 to Mar-03	3.706	3.726	0.055	0.075
One Year Strip*	2.880	2.900	(0.285)	(0.265)

	IF PEPL TX-OK Fixed Price BID	Fixed Price OFFER	Basis BID	Basis OFFER
Cash	2.530	2.550		
ROM	2.530	2.550		
Dec-01	2.828	2.848	(0.133)	(0.113)
Dec-01 to Mar-02	2.958	2.978	(0.130)	(0.110)
Apr-02 to Oct-02	3.046	3.066	(0.120)	(0.100)
Nov-02 to Mar-03	3.531	3.551	(0.120)	(0.100)
One Year Strip*	3.041	3.061	(0.123)	(0.103)

Source: Extract tables from messy spreadsheets with jailbreakr, http://blog.revolutionanalytics.com/2016/08/jailbreakr.html

Data beyond the “raw” stage are written to file

Write functions

```
write_csv()  # use CSV generally
saveRDS()    # use RDS to preserve factors
```

Read functions for further data carpentry

```
read_csv()
readRDS()
```

Data beyond the “raw” stage reside in the `data` directory

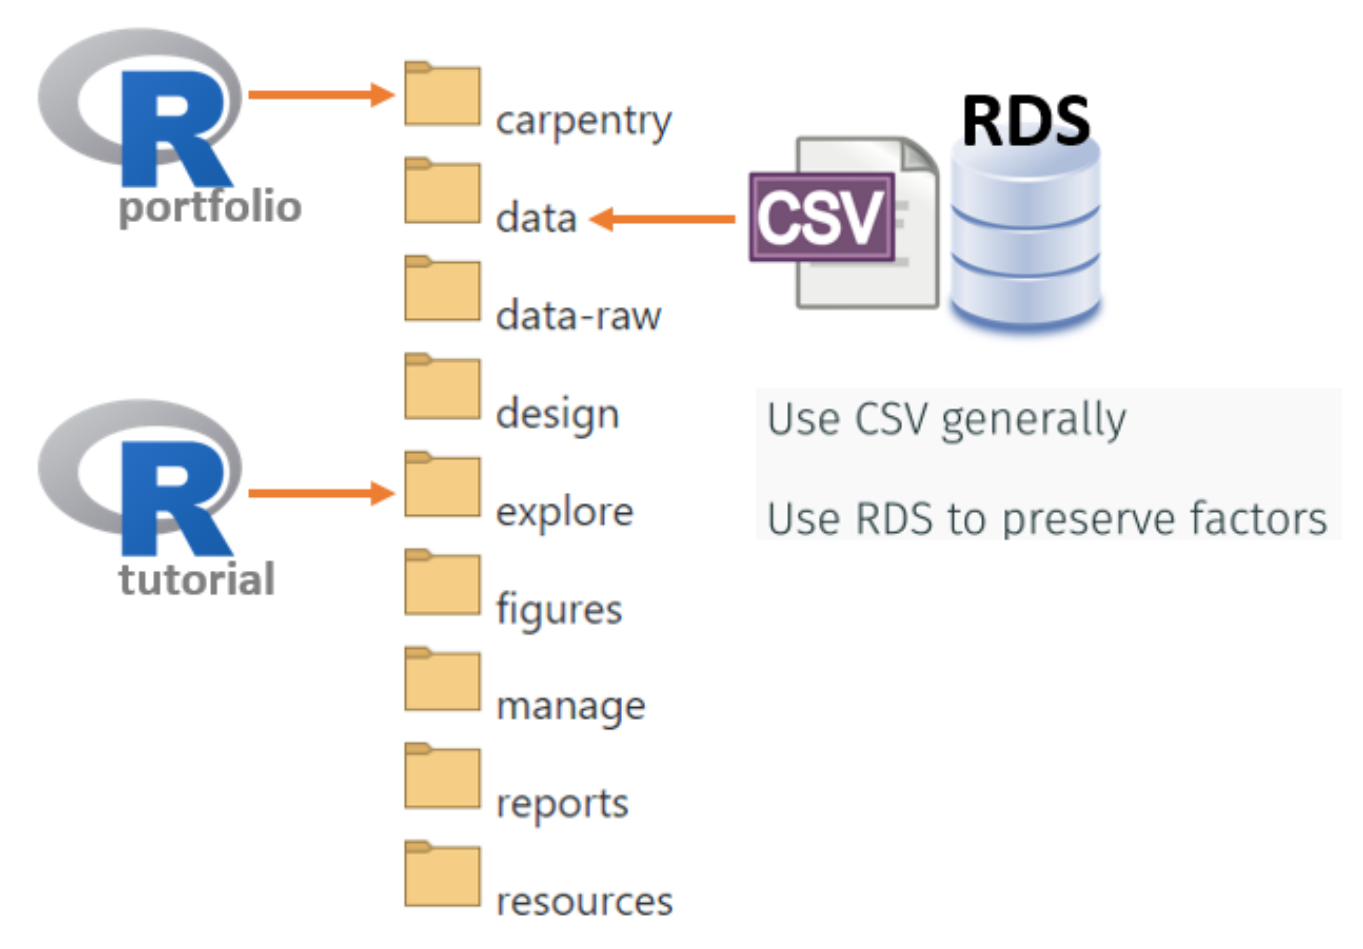

In the data studio, you'll start practicing the skills we've outlined

Obtain the raw data

Read raw data into R and examine it

Identify the structure of your data

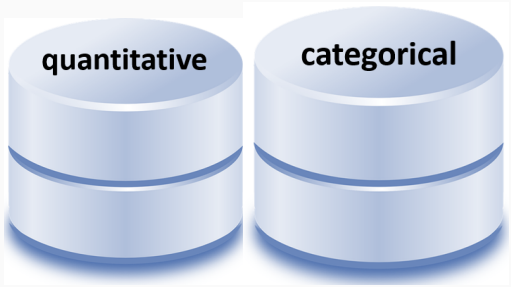

Tidy the data and write to file